

Nonlinear-Equation Solving & Numerical Optimization

Shunsuke Tsuda

ECON 2020 Computing for Economists
Spring 2021



Introduction

- Many economic or econometric problems do not have closed-form solutions.
- Two types of problems to numerically solve:
 - (1) Solving a set of equations
 - (2) Maximizing or minimizing an objective function

Introduction

- Many economic or econometric problems do not have closed-form solutions.
- Two types of problems to numerically solve:
 - (1) Solving a set of equations
 - (2) Maximizing or minimizing an objective function
- Prioritize (1) and avoid (2) whenever possible, because numerical optimization is costly in terms of:
 - (i) Computational time
 - (ii) Risk of missing the exact solution
- However, there are also many situations where numerical optimization is necessary for solving problems.

Motivating Eg. from Applied Microeconomics

- BLP random coefficient logit demand models
- “BLP”: Berry, Levinsohn, and Pakes (1995 ECTA)
“Automobile prices in market equilibrium”
- Influential research in industrial organization and widely imported by other fields (trade, urban, education, health, development, etc)
- Illustrate the estimation of nonlinear models, where the objective function may not be globally concave or convex
- Illustrate the practicality of a wide range of numerical methods (nonlinear equation solving; numerical optimization; numerical differentiation; numerical integration)
- To study BLP, let's first review the homogeneous logit (white board)
- Recommended: [Train “Discrete Choice Methods with Simulation”](#)

BLP RC-logit Demand Model: Setup

- Consumer i 's utility from product j in market t :

$$u_{ijt} = x_j \beta_i - \alpha_i p_{jt} + \xi_{jt} + \varepsilon_{ijt}$$

where

- p_{jt} : price of product j in market t
- x_j : row vec. of non-price characteristics of j
- ξ_{jt} : **product- & market-specific demand shock**
 $\mathbb{E}[\xi_{jt}|p_{jt}, x_j] \neq 0$: **endogeneity** of prices
- **Discrete choice model**: Each consumer chooses one product $j \in 1, \dots, J$ in the market or an outside option $j = 0$ with $u_{i0t} = \varepsilon_{i0t}$, which maximizes his/her utility

BLP RC-logit Demand Model: Setup

- Consumer i 's utility from product j in market t :

$$u_{ijt} = x_j \beta_i - \alpha_i p_{jt} + \xi_{jt} + \varepsilon_{ijt}$$

where

- p_{jt} : price of product j in market t
- x_j : row vec. of non-price characteristics of j
- ξ_{jt} : **product- & market-specific demand shock**
 $\mathbb{E}[\xi_{jt}|p_{jt}, x_j] \neq 0$: **endogeneity** of prices
- **Discrete choice model**: Each consumer chooses one product $j \in 1, \dots, J$ in the market or an outside option $j = 0$ with $u_{i0t} = \varepsilon_{i0t}$, which maximizes his/her utility
- Sources of **consumer heterogeneity**:
 - $\varepsilon_{ijt} \sim_{i.i.d}$ Type I extreme value distribution
 - $(\alpha_i, \beta_i) \sim N(\mu, \Sigma)$

BLP RC-logit Demand Model: Intuition

- Consumer heterogeneity in preferences over product characteristics
⇒ Flexible substitution patterns across products (compared to the simple homogeneous logit)
- Product- & market-specific demand shock unobserved by the econometrician
⇒ Address endogeneity caused by this, employing IVs
- Estimate consumer demand of differentiated products (e.g., cars, cereals) using only widely available aggregate data at the market level

BLP RC-logit Demand Model: Estimation

Step 0 Guess parameter values

Step 1 Define a function that, given ξ_{jt} , predicts market share:

$$\tilde{S}_{jt} = \int \frac{\exp(x_j\beta_i - \alpha_i p_{jt} + \xi_{jt})}{\sum_k \exp(x_k\beta_k - \alpha_k p_{kt} + \xi_{kt})} dF(\alpha, \beta)$$

by **numerical integration**

Step 2 Solve for $\{\xi_{jt}\}$

s.t. Predicted market share = Observed market share

by **nonlinear equation solving**

Step 3 Construct the GMM objective function using instruments

Step 4 Repeat Step 0. – Step 3. for minimizing the GMM objective function by **numerical optimization** to estimate parameter values

Nonlinear Equations

Motivation

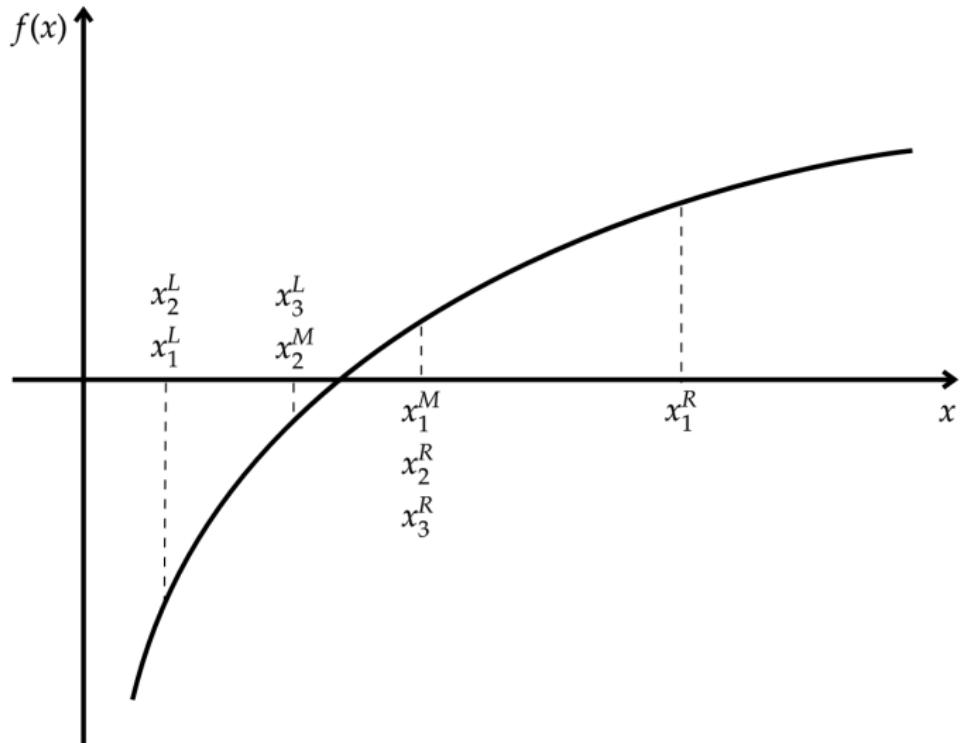
Economic equilibrium is often characterized by systems of nonlinear equations. For example,

- $z(\mathbf{p}) = 0$: System of excess demands with the price vector \mathbf{p} in general equilibrium
- Set of FOCs as conditions for Nash equilibria of games with continuous strategies
(e.g., Cournot competition)
- $k^* = f(k^*)$: Finding a fixed point
(e.g., Steady state of dynamic models; Value function iteration)

Overview

- Bisection method
- Newton's method
- Secant & Broyden's methods
- As a fixed-point problem
- As a minimization problem
- SciPy optimization package ([tutorial](#))
- Further reading: *Judd, Chapter 5*

Bisection Method



Bisection Method: Algorithm

- Based on the Intermediate Value Theorem:
 $f : \mathbb{R} \rightarrow \mathbb{R}$ continuous and $f(a) < 0 < f(b)$
Then, $\exists c \in (a, b)$ s.t. $f(c) = 0$

Bisection Method: Algorithm

- Based on the Intermediate Value Theorem:
 $f : \mathbb{R} \rightarrow \mathbb{R}$ continuous and $f(a) < 0 < f(b)$
Then, $\exists c \in (a, b)$ s.t. $f(c) = 0$

Step 0 Find x^L & x^R s.t. $f(x^L)f(x^R) < 0$
Choose stopping criterion¹: ϵ or δ

Step 1 Compute midpoint: $x^M = (x^L + x^R)/2$

Step 2 Update $x^L = x^L$ & $x^R = x^M$ if $f(x^L)f(x^M) < 0$
Update $x^L = x^M$ & $x^R = x^R$ if $f(x^L)f(x^M) > 0$

Step 3 Check stopping criterion:
If $x^R - x^L \leq \epsilon(1 + |x^L| + |x^R|)$ or $|f(x^M)| \leq \delta$,
stop and report the solution at x^M
Otherwise, go back to Step 1 again

Bisection Method: Pros and Cons

Pros:

- The simplest and most robust method

Bisection Method: Pros and Cons

Pros:

- The simplest and most robust method
- Simplicity: only requires continuity of function
- Robustness: as long as $f(x^L)f(x^R) < 0$, it achieves global convergence with any initial guesses (in case there is only a single root)

Bisection Method: Pros and Cons

Pros:

- The simplest and most robust method
- Simplicity: only requires continuity of function
- Robustness: as long as $f(x^L)f(x^R) < 0$, it achieves global convergence with any initial guesses (in case there is only a single root)

Cons:

- Applicable only to 1-dim rootfinding problems
- Slow: require more iterations than other methods, because it ignores information on the function's curvature

Newton's Method

- Take advantage of information on derivatives of a function
(Recall: Bisection method used only continuity of a function)

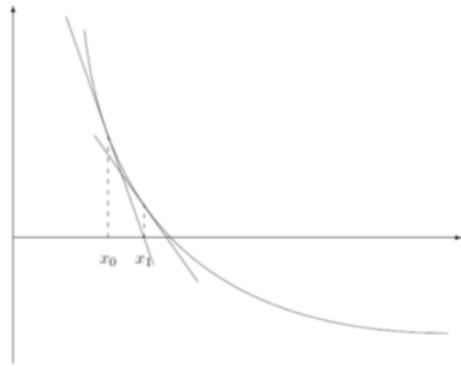
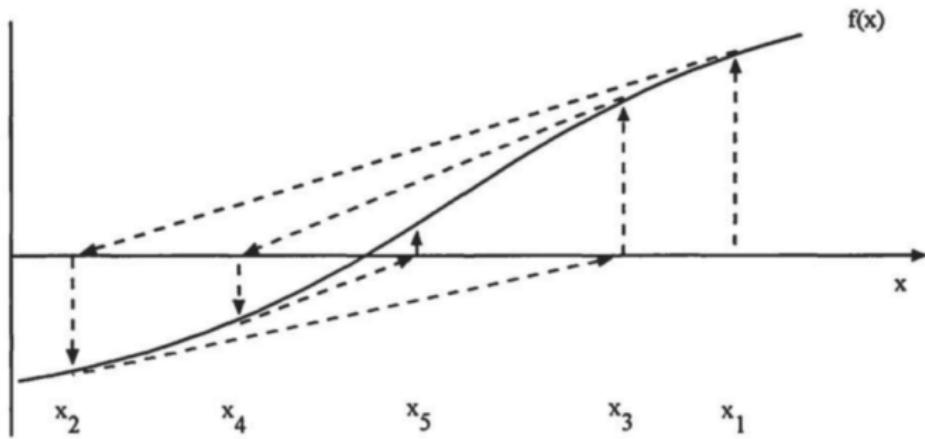
Newton's Method

- Take advantage of information on derivatives of a function
(Recall: Bisection method used only continuity of a function)
- Reduce a nonlinear problem to a sequence of linear problems, where zeros are easy to compute
- $0 = f(x^*) \approx f(x_k) + (x^* - x_k)f'(x_k) \equiv g(x_k)$
- Instead of solving $f(x) = 0$, solve $g(x) = 0$

Newton's Method

- Take advantage of information on derivatives of a function
(Recall: Bisection method used only continuity of a function)
- Reduce a nonlinear problem to a sequence of linear problems, where zeros are easy to compute
- $0 = f(x^*) \approx f(x_k) + (x^* - x_k)f'(x_k) \equiv g(x_k)$
- Instead of solving $f(x) = 0$, solve $g(x) = 0$
- Trade-off: Compared to bisection, faster when it works, but may not always converge

Newton's Method



(Judd, Chapter 5 & Collard's lecture note)

Newton's Method: Algorithm

Step 0 Set x_k with $k = 0$

Choose stopping criteria: ϵ & δ

Step 1 Compute $x_{k+1} = x_k - \frac{f(x_k)}{f'(x_k)}$

Step 2 Check stopping criterion:

If $|x_k - x_{k+1}| \leq \epsilon(1 + |x_{k+1}|)$, go to Step 3.

Otherwise, go back to Step 1.

Step 3 If $|f(x_{k+1})| \leq \delta$, report $x^* = x_{k+1}$ as a solution.

Otherwise, report failure.

Newton's Method: Algorithm

Step 0 Set x_k with $k = 0$

Choose stopping criteria: ϵ & δ

Step 1 Compute $x_{k+1} = x_k - \frac{f(x_k)}{f'(x_k)}$

Step 2 Check stopping criterion:

If $|x_k - x_{k+1}| \leq \epsilon(1 + |x_{k+1}|)$, go to Step 3.

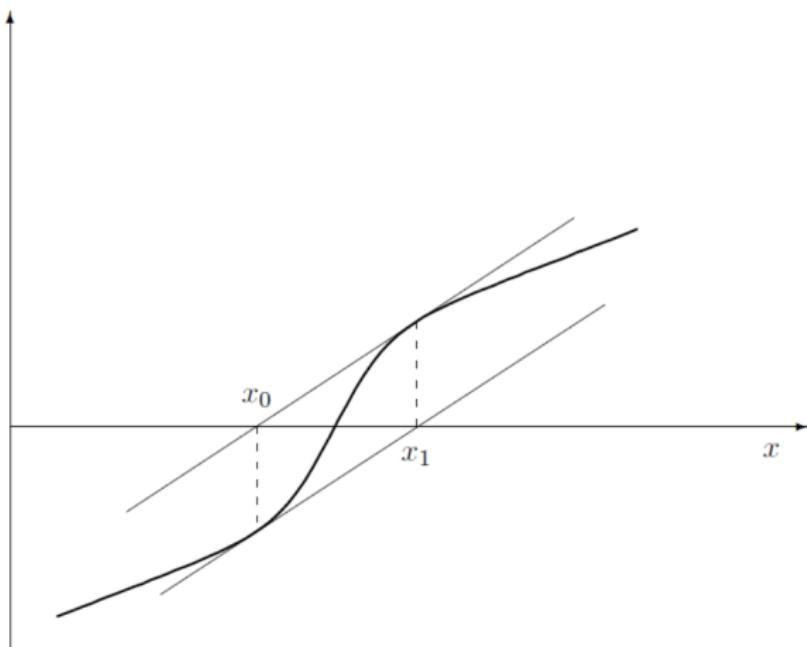
Otherwise, go back to Step 1.

Step 3 If $|f(x_{k+1})| \leq \delta$, report $x^* = x_{k+1}$ as a solution.

Otherwise, report failure.

- Sufficient condition for convergence: x_0 is sufficiently close to x^* , $f'(x^*) \neq 0$, and $|f''(x^*)/f'(x^*)| < \infty$.
(See Theorem 2.1. in *Judd, Chapter 5*)

Remark. x_0 should be *sufficiently* close to x^*



One practical way: First use bisection to obtain a crude approximation for the root, and then shift to Newton.

Newton's Method (Multidimensional)

- $f : \mathbb{R}^n \rightarrow \mathbb{R}^n$; $J_f()$: Jacobian of $f()$
 $\mathbf{0} = f(\mathbf{x}^*) \approx f(\mathbf{x}_k) + J_f(\mathbf{x}_k)(\mathbf{x}^* - \mathbf{x}_k)$

Newton's Method (Multidimensional)

- $f : \mathbb{R}^n \rightarrow \mathbb{R}^n$; $J_f()$: Jacobian of $f()$
 $\mathbf{0} = f(\mathbf{x}^*) \approx f(\mathbf{x}_k) + J_f(\mathbf{x}_k)(\mathbf{x}^* - \mathbf{x}_k)$

Step 0 Set \mathbf{x}_k with $k = 0$

Choose stopping criteria: ϵ & δ

Step 1 Compute $\mathbf{x}_{k+1} = \mathbf{x}_k - J_f(\mathbf{x}_k)^{-1}f(\mathbf{x}_k)$

Step 2 Check stopping criterion:

If $\|\mathbf{x}_k - \mathbf{x}_{k+1}\| \leq \epsilon(1 + \|\mathbf{x}_{k+1}\|)$, go to Step 3.

Otherwise, go back to Step 1.

Step 3 If $\|f(\mathbf{x}_{k+1})\| \leq \delta$, report $\mathbf{x}^* = \mathbf{x}_{k+1}$ as a solution.
Otherwise, report failure.

Remark. Require $\det(J(x^*)) \neq 0$

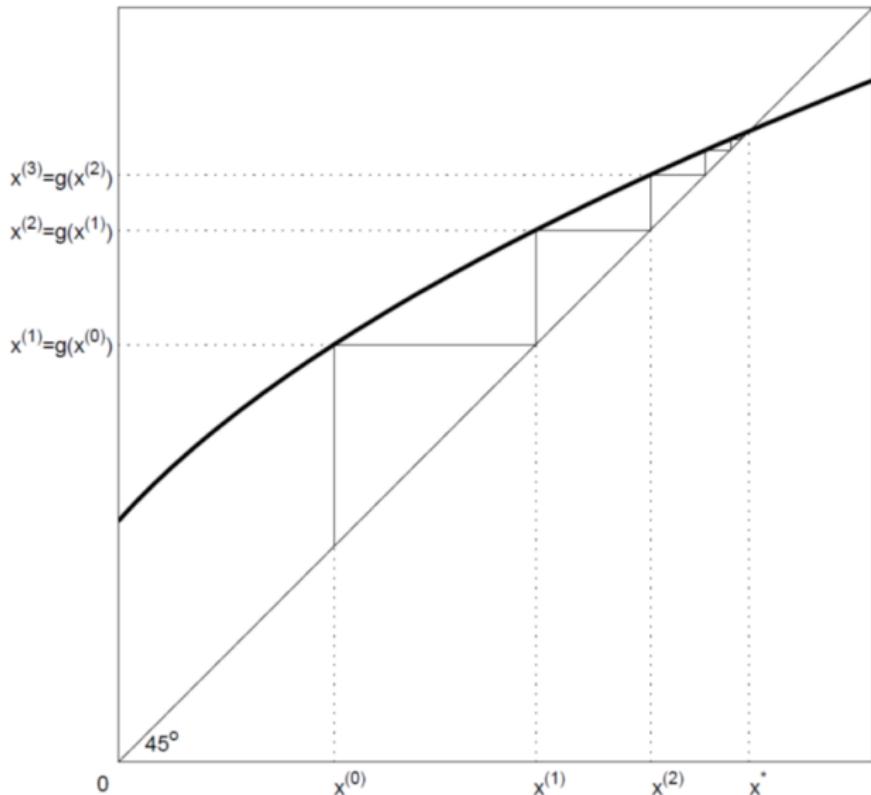
Secant Method

- Obtaining $f'(x_k)$ or Jacobian is costly to compute and code
- Secant method instead uses the simplest approximations of $f'(x_k)$ & $J_f(\mathbf{x}_k)$ by m_k & A_k s.t.
$$f(x_k) - f(x_{k-1}) = m_k(x_k - x_{k-1}) \text{ for one-dimensional}$$
$$f(\mathbf{x}_k) - f(\mathbf{x}_{k-1}) = A_k(\mathbf{x}_k - \mathbf{x}_{k-1}) \text{ for multi-dimensional}$$
- Called Broyden's method in a multidimensional case
- Set (x_0, x_1) at the beginning
- Other processes and properties are similar to Newton's method
- See *Judd, Chapter 5* in detail

As a Fixed-Point Problem

- Fixed-point problems arise frequently in economic problems.
- Any fixed point problems: $x = g(x)$ can be cast as a nonlinear-eq solving: $f(x) \equiv x - g(x) = 0$
- Some (not all) nonlinear-eq solving problems can be cast as a fixed point problem

Fixed-Point Iteration



(Miranda and Fackler, Chapter 3)

Fixed-Point Iteration

- Compute a fixed point of $x = g(x)$ ($g : \mathbb{R} \rightarrow \mathbb{R}$)
- Construct a sequence $\{x_k\}$ s.t. $x_{k+1} = g(x_k)$
- If it converges ($\lim_{k \rightarrow \infty} x_k = x^*$), then $x^* = g(x^*)$

Fixed-Point Iteration

- Compute a fixed point of $x = g(x)$ ($g : \mathbb{R} \rightarrow \mathbb{R}$)
- Construct a sequence $\{x_k\}$ s.t. $x_{k+1} = g(x_k)$
- If it converges ($\lim_{k \rightarrow \infty} x_k = x^*$), then $x^* = g(x^*)$
- *Contraction mapping theorem:* If there exists $T \in [0, 1)$ such that $|g(x) - g(x')| \leq T|x - x'| \quad \forall x, x' \in D \subset \mathbb{R}$, then x^* is a unique fixed point in D .

Fixed-Point Iteration

- Compute a fixed point of $x = g(x)$ ($g : \mathbb{R} \rightarrow \mathbb{R}$)
- Construct a sequence $\{x_k\}$ s.t. $x_{k+1} = g(x_k)$
- If it converges ($\lim_{k \rightarrow \infty} x_k = x^*$), then $x^* = g(x^*)$
- *Contraction mapping theorem:* If there exists $T \in [0, 1)$ such that $|g(x) - g(x')| \leq T|x - x'| \quad \forall x, x' \in D \subset \mathbb{R}$, then x^* is a unique fixed point in D .
- Finding such D is hard. In that case, a modified updating scheme with *extrapolation* is preferable to stabilize:

$$x_{k+1} = \lambda_k x_k + (1 - \lambda_k)g(x_k)$$

where $\lambda_k \in [0, 1]$ and $\lim_{k \rightarrow \infty} \lambda_k = 0$

Fixed-Point Iteration

- Compute a fixed point of $x = g(x)$ ($g : \mathbb{R} \rightarrow \mathbb{R}$)
- Construct a sequence $\{x_k\}$ s.t. $x_{k+1} = g(x_k)$
- If it converges ($\lim_{k \rightarrow \infty} x_k = x^*$), then $x^* = g(x^*)$
- *Contraction mapping theorem:* If there exists $T \in [0, 1)$ such that $|g(x) - g(x')| \leq T|x - x'| \quad \forall x, x' \in D \subset \mathbb{R}$, then x^* is a unique fixed point in D .
- Finding such D is hard. In that case, a modified updating scheme with *extrapolation* is preferable to stabilize:

$$x_{k+1} = \lambda_k x_k + (1 - \lambda_k)g(x_k)$$

where $\lambda_k \in [0, 1]$ and $\lim_{k \rightarrow \infty} \lambda_k = 0$

- Trade-off between accuracy and speed exists
- On the other hand, if the original system is converging too slowly, $\lambda_k < 0$ could be a way to accelerate convergence

Fixed-Point Iteration

```
import matplotlib.pyplot as plt
import numpy as np
from numpy.linalg import norm
from scipy.optimize import bisect, newton

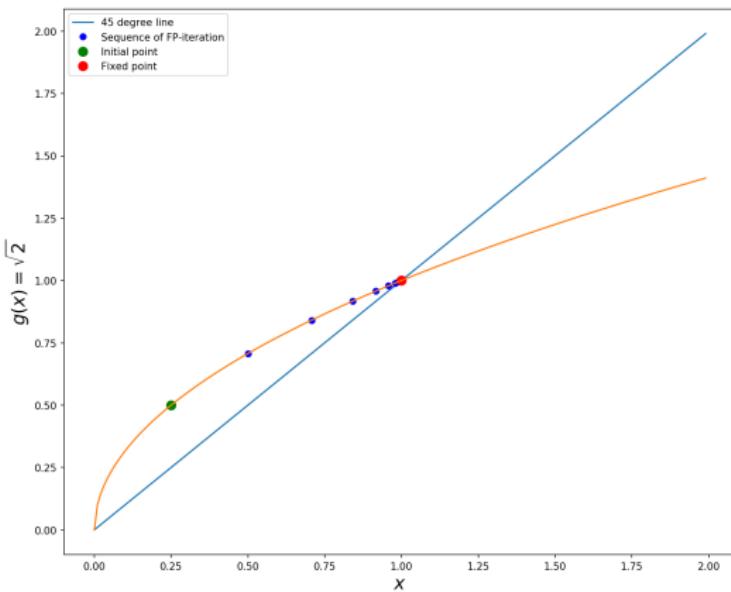
def fp_iter(g,x0,tol=10e-8,maxiter=100):
    """ Fixed point iteration """
    e = 1 # error
    iter = 0 # number of iteration
    x_seq = [] # store the sequence
    while(e > tol and iter < maxiter):
        ###You will code by yourself in the lab session###
    return x,x_seq

def fp_iter_rev(g,x0,lambda_k,tol=10e-8,maxiter=100):
    """ Fixed point iteration with alternative updating scheme"""
    ###You will code by yourself in the lab session###
return x,x_seq

g1 = lambda x : np.sqrt(x)      # E.g.1)
x_init = .25
x1_fp, x1_seq = fp_iter(g1,x_init)

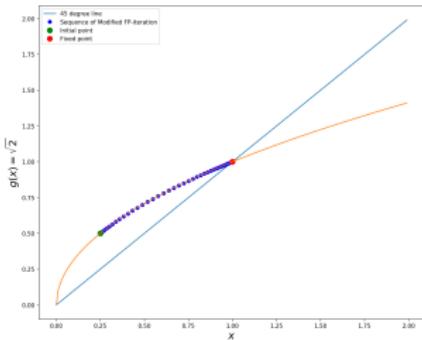
g2 = lambda x : x**2 - 1      # E.g.2)
x_init = -0.5
lambda_init = 0.99
x2_fp, x2_seq = fp_iter(g2,x_init)
x2_fp_rev, x2_seq_rev = fp_iter_rev(g2,x_init,lambda_init)
```

E.g.1) Convergence by the Benchmark Scheme



```
In [1]: x1_seq
Out[1]:
array([0.5      , 0.70710678, 0.84089642, 0.91700404, 0.95760328,
       0.97857206, 0.98922801, 0.99459942, 0.99729606, 0.99864711,
       0.99932333, 0.99966161, 0.99983079, 0.99991539, 0.99995769,
       0.99997885, 0.99998942, 0.99999471, 0.99999736, 0.99999868,
       0.99999934, 0.99999967, 0.99999983, 0.99999992])
```

E.g.1) Slower Convergence by the Modified Scheme

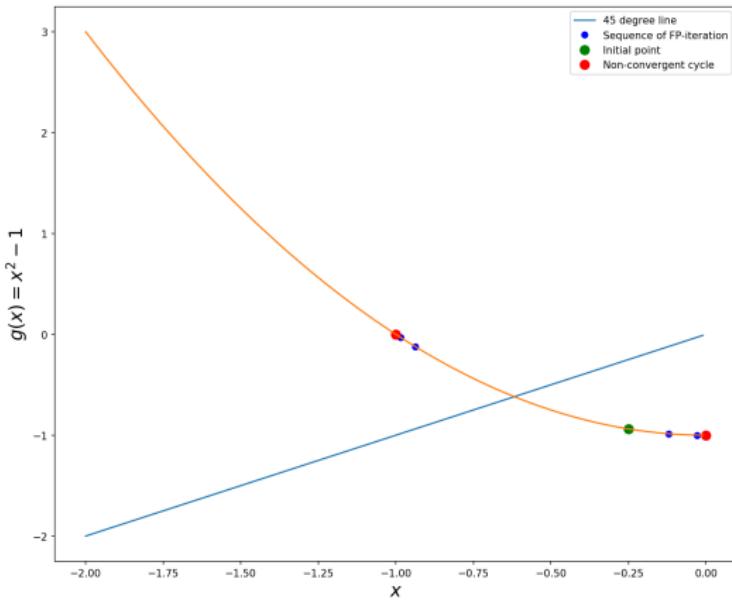


```
In [1]: x1_seq_rev
```

```
Out[1]:
```

```
array([0.2525      , 0.25747488, 0.26489849, 0.27474099, 0.28696488,
       0.3015203   , 0.31834012, 0.33733537, 0.35839136, 0.38136516,
       0.40608443  , 0.43234806, 0.45992833, 0.4885748  , 0.51801943,
       0.54798281  , 0.57818103, 0.60833281, 0.63816647, 0.66742641,
       0.69587881  , 0.72331621, 0.749561  , 0.77446754, 0.79792312,
       0.81984774  , 0.84019279, 0.85893899, 0.87609352, 0.89168681,
       0.90576896  , 0.91840611, 0.92967684, 0.93966874, 0.94847527,
       0.95619295  , 0.96291894, 0.96874895, 0.97377564, 0.9780873  ,
       0.98176696  , 0.98489171, 0.98753235, 0.98975326, 0.99161242,
       0.99316159  , 0.99444663, 0.99550781, 0.9963803  , 0.99709453,
       0.99767672  , 0.99814927, 0.99853123, 0.9988387  , 0.9990852  ,
       0.99928202  , 0.99943855, 0.99956254, 0.99966037, 0.99973726,
       0.99979747  , 0.99984442, 0.99988091, 0.99990916, 0.99993095,
       0.99994769  , 0.9999605  , 0.99997028, 0.99997771, 0.99998334,
       0.99998759  , 0.99999079, 0.99999318, 0.99999497, 0.9999963  ,
       0.99999729  , 0.99999802, 0.99999856, 0.99999895, 0.99999924,
       0.99999945  , 0.99999961, 0.99999972, 0.9999998  , 0.99999986,
       0.9999999 ])
```

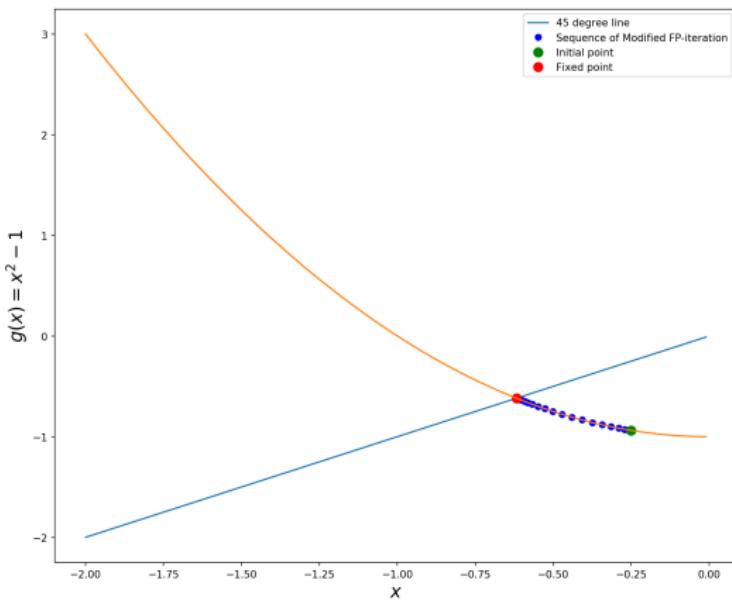
E.g.2) Convergence Failure by the Benchmark Scheme



In [1]: x2_seq

Out[2]:

E.g.2) Convergence by the Modified Scheme



```
In [1]: x2_seq_rev
```

```
Out[1]:
```

```
array([-0.256875, -0.27035009, -0.2898506, -0.31452286, -0.34326977,
       -0.37480589, -0.40773481, -0.44064692, -0.47222899, -0.50137054,
       -0.52724889, -0.54937654, -0.5676025, -0.58206998, -0.5931428,
       -0.60131839, -0.60714485, -0.61115463, -0.6138204, -0.61553283,
       -0.61659578, -0.61723331, -0.61760269, -0.61780934, -0.61792094,
       -0.61797907, -0.61800825, -0.61802237, -0.61802894, -0.61803188,
       -0.61803314, -0.61803366, -0.61803387, -0.61803395, -0.61803397])
```

Fixed-Point Iteration (Multidimensional)

- Fixed point problem in a multidimensional system with $g : \mathbb{R}^n \rightarrow \mathbb{R}^n$

$$x_1 = g_1(x_1, \dots, x_n)$$

$$\vdots$$

$$x_n = g_n(x_1, \dots, x_n)$$

- Construct a sequence $\{\mathbf{x}_k\}$ s.t. $\mathbf{x}_{k+1} = g(\mathbf{x}_k)$
- Contraction mapping theorem applies.
(Check the properties related to Jacobian.)
- The modified updating scheme can also be used similarly to the single-dimensional case

Anderson Acceleration for Fixed-Point Iterations

- Alleviates the potential concern of slow convergence or divergence associated with the standard FP iteration
- This nests the standard FP iteration, i.e., equivalent under some parameter values
- See Walker and Ni (2011) for detail
- Both Python (Scipy) and Julia (NLsolve.jl) contain the package to implement the Anderson Acceleration

Other Elementary Methods for a Multidimensional System

Instead of solving n equations for n unknowns, repeatedly solve each one of n equations with one unknown in turn:

- Gauss-Jacobi Algorithm
- Gauss-Seidel Algorithm

Toward Convergence

- The convergence is not guaranteed in most methods for a multi-dimensional system
- Sensitive to initial guesses, but often difficult to find good ones

Toward Convergence

- The convergence is not guaranteed in most methods for a multi-dimensional system
- Sensitive to initial guesses, but often difficult to find good ones
- Obtain better initial guesses from **optimization** ideas:
 - Optimization problems are less sensitive to initial guesses
 - One simple way is to obtain a rough guess of $f(\mathbf{x}) = 0$ by solving $\min_{\mathbf{x}} \sum_{i=1}^n f^i(\mathbf{x})^2$ with a loose stopping rule
 - (We cover optimization algorithms in the next section)

Toward Convergence

- The convergence is not guaranteed in most methods for a multi-dimensional system
- Sensitive to initial guesses, but often difficult to find good ones
- Obtain better initial guesses from **optimization** ideas:
 - Optimization problems are less sensitive to initial guesses
 - One simple way is to obtain a rough guess of $f(\mathbf{x}) = 0$ by solving $\min_{\mathbf{x}} \sum_{i=1}^n f^i(\mathbf{x})^2$ with a loose stopping rule
 - (We cover optimization algorithms in the next section)
- **Continuation methods:** Construct a sequence of problems (each of which is reasonably solvable) that ultimately leads to the problem of interest

Lab

Exercise

Aim. Get accustomed to basic one-dimensional methods for solving a non-linear equation

- $f(x) \equiv \exp((x - 2)^2) - 2 - x$
 $g(x) \equiv \exp((x - 2)^2) - 2$
- Solve $f(x) = 0$ or $x = g(x)$ over the domain $x \in [0, 2]$

Tasks.

1. Solve $f(x) = 0$ by the bisection method with initial $x^L = 0.5$ & $x^R = 1.5$.
You can use the [SciPy root-finding package](#).
 - 2.1. Solve $x = g(x)$ by the fixed-point iteration with the updating rule $x_{k+1} = g(x_k)$. Code the algorithm by yourself.
Does this work? Explain why or why not.
 - 2.2. Solve $x = g(x)$ by the fixed-point iteration with the updating rule $x_{k+1} = \lambda_k x_k + (1 - \lambda_k)g(x_k)$ where $\lambda_0 = 1$ & $\lambda_k = 0.99\lambda_{k-1}$
- Using `matplotlib.pyplot`, plot the convergent sequence (like what we saw in the lecture).

(Adapted from *Collard, Lecture Notes 5.*)

Assignment 3

Assignment: Cournot Duopoly Model

Aim. Get accustomed to basic multi-dimensional methods for solving a system of non-linear equations. Vectorize the system.

- Quantity competition by 2 firms $i = 1, 2$
- Inverse demand of a good: $P(q) = q^{-1/\alpha}$
- Cost function: $C_i(q_i) = \frac{1}{2}c_i q_i^2$
- Profit for each firm i : $\pi_i(q_1, q_2) = P(q_1 + q_2)q_i - C_i(q_i)$
- Assume $c_1 = 0.6$ & $c_2 = 0.8$

Task 1. Solve for equilibrium q_i^* for $i = 1, 2$ and $P(q_1^* + q_2^*)$ with $\alpha = 1.5$ by

- (1) Newton's method: Code the algorithm by yourself. That is, analytically derive the Jacobian of the set of FOCs.
- (2) Broyden's method: You can use the [SciPy root-finding package](#).
- (3) Fixed-point iteration: Code the algorithm by yourself.

Task 2. Solve the model for all $\alpha \in [1, 3]$ (construct grids) by Broyden's method. Using `matplotlib.pyplot`, produce two plots with x -axis α & y -axes (i) q_1^*, q_2^* (ii) $P(q_1^* + q_2^*)$.

Note for Non-Economics Students

- **Imperfect competition:** There are only two firms in the market. Each firm's quantity decision impacts the market price via consumer demand structure.
- The first order condition of firm i 's profit maximization problem (given the other firm's decision):

$$\frac{\partial \pi_i}{\partial q_i} = P(q_1 + q_2) + P'(q_1 + q_2)q_i - C'_i(q_i) = 0$$

- The equilibrium condition is that both the following equalities hold simultaneously:

$$f_1(q_1, q_2) \equiv (q_1 + q_2)^{-\frac{1}{\alpha}} - \frac{1}{\alpha}(q_1 + q_2)^{-\frac{1}{\alpha}-1}q_1 - c_1q_1 = 0$$

$$f_2(q_1, q_2) \equiv (q_1 + q_2)^{-\frac{1}{\alpha}} - \frac{1}{\alpha}(q_1 + q_2)^{-\frac{1}{\alpha}-1}q_2 - c_2q_2 = 0$$

- Numerically solve for the equilibrium q_1^* & q_2^* .

Optimization

Numerical Optimization: Motivation

- Optimization is ubiquitous in economics and econometrics
- **Economic problems:** Consumer's utility maximization, Firm's profit maximization and cost minimization, Social planner's total surplus maximization, etc
- **Econometric problems:** Minimizing the sum of squared errors, Minimizing the GMM objective function, Maximizing the likelihood function, etc

Numerical Optimization: Remarks

- Numerical optimization is costly in terms of:
 - (i) Computation time
 - (ii) Risk of missing the exact solution

Hence, avoid it whenever possible

Numerical Optimization: Remarks

- Numerical optimization is costly in terms of:
 - (i) Computation time
 - (ii) Risk of missing the exact solution

Hence, avoid it whenever possible

- (i): Check if solving a set of equations suffices to solve the whole problem, which saves computation time. E.g.:
- Convex optimization problems in which KKT conditions are sufficient for optimization
 - Exactly identified case in GMM

Numerical Optimization: Remarks

- Numerical optimization is costly in terms of:
 - (i) Computation time
 - (ii) Risk of missing the exact solution

Hence, avoid it whenever possible

- (i): Check if solving a set of equations suffices to solve the whole problem, which saves computation time. E.g.:
 - Convex optimization problems in which KKT conditions are sufficient for optimization
 - Exactly identified case in GMM
- Choosing just a faster optimization algorithm is also dangerous due to (ii).
- Important to understand the trade-off between accuracy and speed across different optimization algorithms.

Remarks

- Optimization is costly in terms of
 - (i) Computation time
 - (ii) Risk of missing the exact solution

Hence, avoid it whenever possible

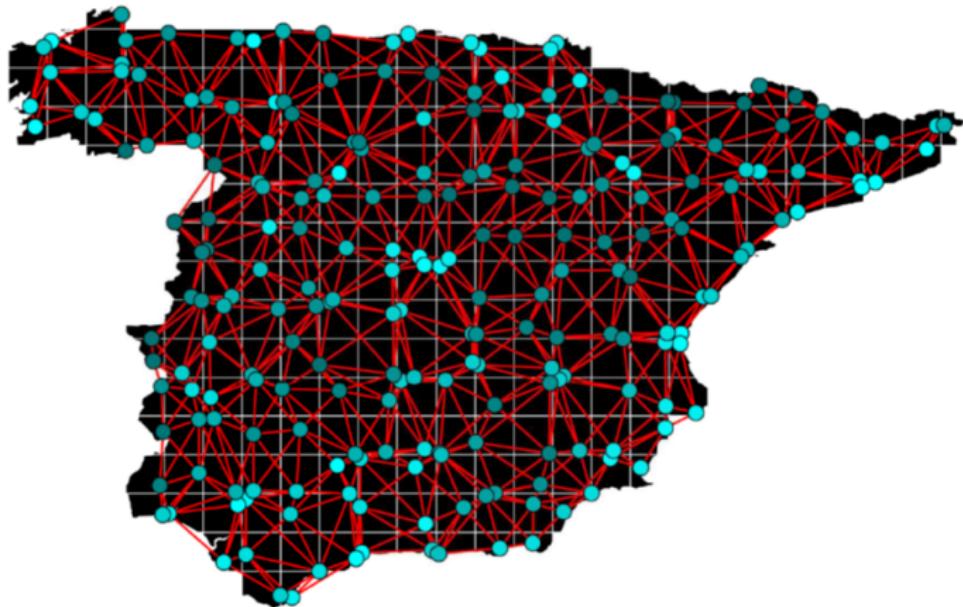
- (ii): Any optimization algorithm finds *local* optimum, but there is no guarantee that *global* optimum is found (unless the objective function is globally concave or convex).

Obtained solutions are susceptible to:

- Search algorithms
- Initial guesses
- Stopping rules

Example of (i): Fajgelbaum & Schaal (2020 ECTA)

- Fajgelbaum & Schaal (2020 ECTA) “Optimal Transport Networks in Spatial Equilibrium”
- Q. How large are the gains from expansion and the losses from misallocation of current road networks in Europe?
- An example of an economically simple, but a computationally hard problem
- Also, an example of research in which quantification is of first order importance
- Illustrate the importance of caring about computational burden in a large-scale problem



The problem of optimally designing the road network is determining how much to invest in each link.

Fajgelbaum & Schaal (2020): Environment

Geography:

- $\mathcal{J} = \{1, \dots, J\}$: locations (nodes)
- $\mathcal{N}(j)$: set of connected location of j
Goods can be directly shipped only through connected locations
- L_j : Number of workers in j , immobile across locations
 L : Total number of workers

Fajgelbaum & Schaal (2020): Environment

Geography:

- $\mathcal{J} = \{1, \dots, J\}$: locations (nodes)
- $\mathcal{N}(j)$: set of connected location of j
Goods can be directly shipped only through connected locations
- L_j : Number of workers in j , immobile across locations
 L : Total number of workers

Commodities:

- $Y_j^n = z_j^n L_j^n$: tradable good production for sector n
- H_j : the non-tradable good endowment (constant)
 $h_j = H_j / L_j$: per-capita consumption of the non-traded good
- $C_j = \left(\sum_{n=1}^N (C_j^n)^{\frac{\sigma-1}{\sigma}} \right)^{\frac{\sigma}{\sigma-1}}$
 $c_j = C_j / L_j$: per-capita consumption of traded goods bundle

Fajgelbaum & Schaal (2020): Environment

Geography:

- $\mathcal{J} = \{1, \dots, J\}$: locations (nodes)
- $\mathcal{N}(j)$: set of connected location of j
Goods can be directly shipped only through connected locations
- L_j : Number of workers in j , immobile across locations
 L : Total number of workers

Commodities:

- $Y_j^n = z_j^n L_j^n$: tradable good production for sector n
- H_j : the non-tradable good endowment (constant)
 $h_j = H_j / L_j$: per-capita consumption of the non-traded good
- $C_j = \left(\sum_{n=1}^N (C_j^n)^{\frac{\sigma-1}{\sigma}} \right)^{\frac{\sigma}{\sigma-1}}$
 $c_j = C_j / L_j$: per-capita consumption of traded goods bundle

Preference:

- $U(c, h)$: worker's utility (homothetic and concave)

Network Building & Transport Technologies

- I_{jk} : Infrastructure (roads) along the link jk
Def “Transport network” = Distribution of $\{I_{jk}\}_{j \in \mathcal{J}, k \in N(j)}$
- K: total resource for building infrastructure
Building I_{jk} requires investment $\delta_{jk}^I I_{jk}$ units of K

Network Building & Transport Technologies

- I_{jk} : Infrastructure (roads) along the link jk
Def “Transport network” = Distribution of $\{I_{jk}\}_{j \in \mathcal{J}, k \in N(j)}$
- K : total resource for building infrastructure
Building I_{jk} requires investment $\delta_{jk}^I I_{jk}$ units of K
- Q_{jk}^n : Quantity of good n shipped from j to $k \in N(j)$
- Transporting one unit of good n from j to $k \in N(j)$ requires τ_{jk}^n units of good n , i.e., “Iceberg cost” = $1 + \tau_{jk}^n$, where

$$\tau_{jk}^n = \tau_{jk}(Q_{jk}^n, I_{jk}) = \delta_{jk}^\tau \frac{(Q_{jk}^n)^\beta}{(I_{jk})^\gamma}$$

δ_{jk}^τ : geographic frictions

(e.g., distance, elevation, ruggedness, river, etc)

$\beta, \gamma > 0$: decreasing returns to transport (congestion force) & positive returns to infrastructure

Network Building & Transport Technologies

- I_{jk} : Infrastructure (roads) along the link jk
Def “Transport network” = Distribution of $\{I_{jk}\}_{j \in \mathcal{J}, k \in N(j)}$
- K : total resource for building infrastructure
Building I_{jk} requires investment $\delta_{jk}^I I_{jk}$ units of K
- Q_{jk}^n : Quantity of good n shipped from j to $k \in N(j)$
- Transporting one unit of good n from j to $k \in N(j)$ requires τ_{jk}^n units of good n , i.e., “Iceberg cost” = $1 + \tau_{jk}^n$, where

$$\tau_{jk}^n = \tau_{jk}(Q_{jk}^n, I_{jk}) = \delta_{jk}^\tau \frac{(Q_{jk}^n)^\beta}{(I_{jk})^\gamma}$$

δ_{jk}^τ : geographic frictions

(e.g., distance, elevation, ruggedness, river, etc)

$\beta, \gamma > 0$: decreasing returns to transport (congestion force) & positive returns to infrastructure

- Total transport costs $Q_{jk}^n \tau_{jk}(Q_{jk}^n, I_{jk})$ are jointly convex over Q_{jk}^n and I_{jk} iff $\beta \geq \gamma$

Fajgelbaum and Schaal (2020): Social Planner's Problem

SP's problem consists of three subproblems:

- Optimal Allocation (given infrastructure and goods flow):
- Optimal Transport (given infrastructure)
- Optimal Infrastructure Network Design
(= the full problem)

$$\begin{aligned}
W &= \max_{\{c_j, h_j, \{I_{jk}\}_{k \in N(j)}, C_j^n, L_j^n, \{Q_{jk}^n\}_{k \in N(j)}\}} \sum_j \omega_j L_j U(c_j, h_j) \\
&= \underbrace{\max_{\{I_{jk}\}_{k \in N(j)}} \max_{\{Q_{jk}^n\}_{k \in N(j)}} \underbrace{\max_{\{c_j, h_j, C_j^n, L_j^n\}} \sum_j \omega_j L_j U(c_j, h_j)}_{\text{Optimal allocation subproblem}}}_{\text{Optimal transport subproblem}} \\
&\quad \underbrace{\qquad\qquad\qquad}_{\text{Optimal infrastructure network design problem}}
\end{aligned}$$

s.t. (i) Availability of tradable & non-tradable goods:

$$c_j L_j \leq C_j \text{ & } h_j L_j \leq H_j \quad \forall j$$

(ii) Balanced-flows constraint:

$$\begin{array}{ccccc}
\underbrace{C_j^n}_{\text{Consumption}} & + & \underbrace{\sum_{k \in N(j)} (1 + \tau_{jk}(Q_{jk}^n, I_{jk})) Q_{jk}^n}_{\text{Exports}} & \leq & \underbrace{Y_j^n}_{\text{Production}} + \underbrace{\sum_{i \in N(j)} Q_{ij}^n}_{\text{Imports}} \quad \forall n, j
\end{array}$$

(iii) Network-building constraint:

$$\sum_j \sum_{k \in N(j)} \delta_{jk}^I I_{jk} \leq K$$

(iv) Local labor market clearing

(v) Non-negativity constraints on consumption, flows, and factor use

Fajgelbaum & Schaal (2020): Social Planner's Problem

SP's problem consists of three subproblems:

- Optimal Allocation (given infrastructure and goods flow)
- Optimal Transport (given infrastructure)
- Optimal Infrastructure Network Design
(= the full problem)

The full problem is **globally convex** if the transport costs are jointly convex, i.e., if $\beta \geq \gamma$
(review the 1st year math!)

Fajgelbaum & Schaal (2020): Numerical Implementation

Convex cases ($\beta \geq \gamma$):

- KKT conditions are both necessary and sufficient
- Numerically tractable

Fajgelbaum & Schaal (2020): Numerical Implementation

Convex cases ($\beta \geq \gamma$):

- KKT conditions are both necessary and sufficient
- Numerically tractable

Non-Convex cases ($\beta < \gamma$):

- The above approach is not guaranteed to find the global optimum
- Optimal transport and allocation subproblems are convex if $Q\tau_{jk}(Q, I_{jk})$ is convex in Q , i.e., if $\beta \geq 0$
- Iterative procedure over the infrastructure investments:
 - Guess on the network investment I_{jk}
 - Solve for the optimum over $\{c_j, C_j^n, h_j, L_j^n, Q_{jk}^n\}$
 - Obtain a new guess over I_{jk}
 - Repeat until convergence...

Example of (ii): BLP Again

- Consumer i 's utility from product j in market t :

$$u_{ijt} = x_j \beta_i - \alpha_i p_{jt} + \xi_{jt} + \epsilon_{ijt}$$

where

- p_{jt} : price of product j in market t
- x_j : row vec. of non-price characteristics of j
- ξ_{jt} : **product- & market-specific demand shock**
 $\mathbb{E}(\xi_{jt}|p_{jt}, x_j) \neq 0$: **endogeneity** of prices
- **Discrete choice model**: Each consumer chooses one product $j \in 1, \dots, J$ in the market or an outside option $j = 0$ with $u_{i0t} = \epsilon_{i0t}$, which maximizes his/her utilit
- Sources of **consumer heterogeneity**:
 - $\epsilon_{ijt} \sim_{i.i.d}$ Type I extreme value distribution
 - $(\alpha_i, \beta_i) \sim N(\mu, \Sigma)$

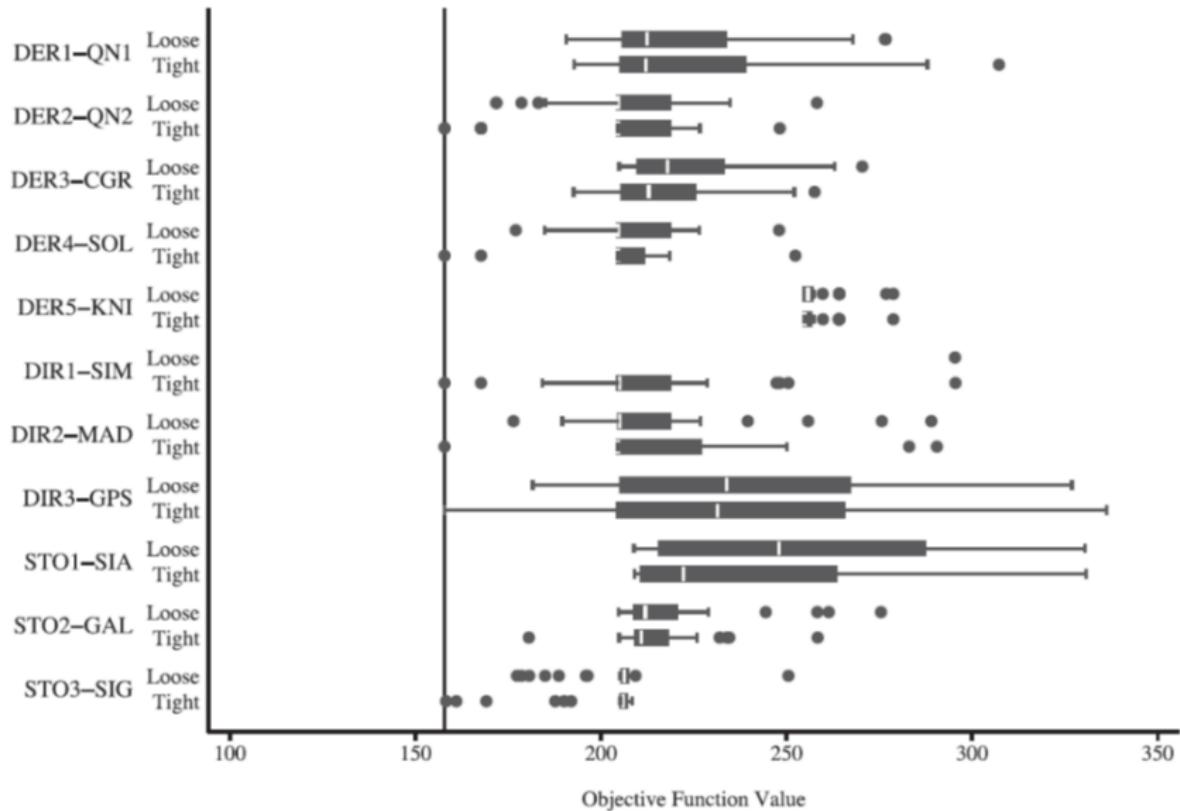
Example of (ii): BLP Again

- Here I follow [Knittel and Metaxoglou \(2014 REStat\)](#) which emphasize numerical challenges of BLP RC-logit demand models
- Illustrates with the BLP RC-logit demand model that different combinations of search algorithms, initial guesses, and stopping rules lead to convergences at different optima
- Observed convergences at:
 - Points where 1st- and 2nd- order conditions fail!
 - Local optima
- Also, observed convergence failure in some instances

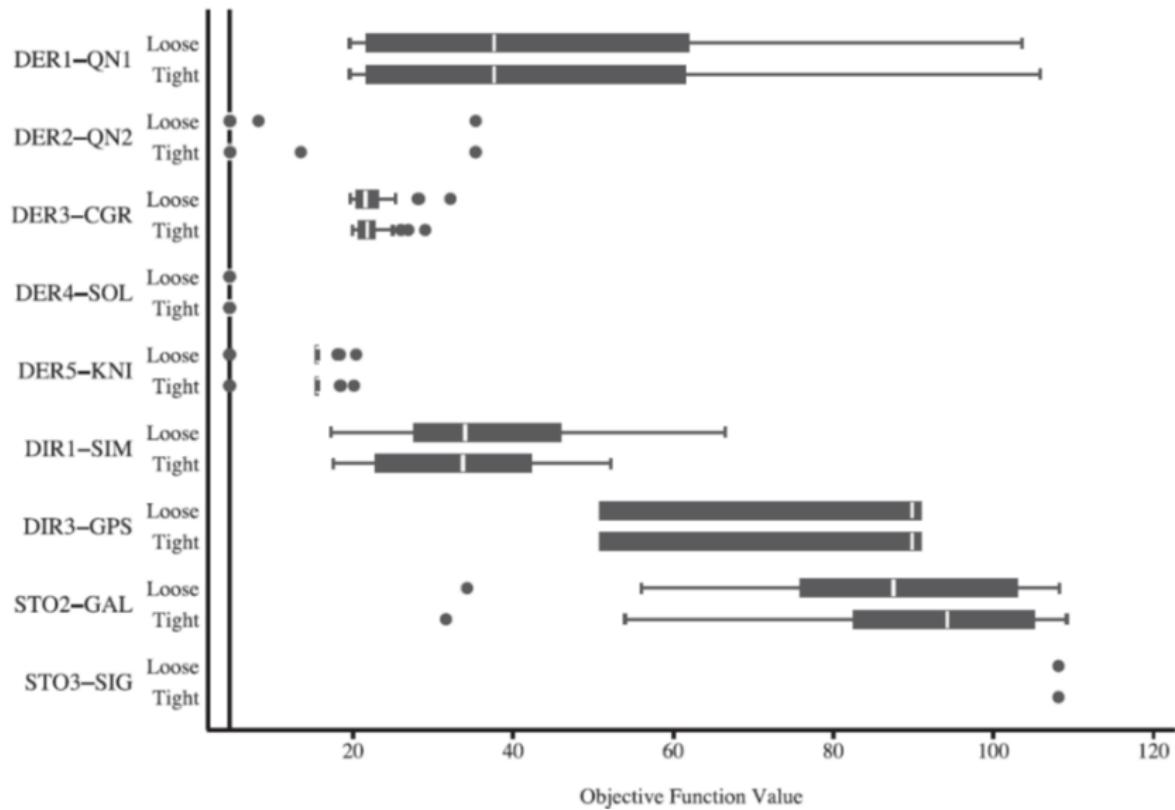
TABLE 2.—OPTIMIZATION ALGORITHMS

Class	Description	Source	Acronym
Derivative-based	Quasi-Newton 1	MathWorks	DER1-QN1
	Quasi-Newton 2	Publicly available	DER2-QN2
	Conjugate gradient	Publicly available	DER3-CGR
	SOLVOPT	Publicly available	DER4-SOL
	KNITRO	Ziena Optimization	DER5-KNI
Deterministic direct search	Simplex	MathWorks	DIR1-SIM
	Mesh adaptive direct search	MathWorks	DIR2-MAD
	Generalized pattern search	MathWorks	DIR3-GPS
Stochastic direct search	Simulated annealing	Publicly Available	STO1-SIA
	Genetic algorithm GADS	MathWorks	STO2-GAL
	Simulated annealing GADS	MathWorks	STO3-SIG

A. Automobiles



B. Cereals



Summary

- Optimization is costly in terms of
 - (i) Computation time
 - (ii) Risk of missing the exact solution

Hence, avoid it whenever possible

- (i): Check if solving a set of equations suffices to solve the whole problem, which saves computation time.
- (ii): Any optimization algorithm finds *local* optimum, but there is no guarantee that *global* optimum is found (unless the objective function is globally concave or convex).
- Understand the trade-off between accuracy and speed across different optimization algorithms.

3 Types of Methods

- Derivative-based methods (1st order):
 - Require differentiability
 - Uses information about gradients
- Derivative-based methods (2nd order):
 - Uses information about gradients and curvature
 - Converges more rapidly to the solution
 - High cost of computing and storing Hessians
- Derivative-free methods (0th order):
 - Slow
 - Suitable for problems with kinks and discontinuities

3 Types of Methods

- Derivative-based methods (1st order):
 - Require differentiability
 - Uses information about gradients
- Derivative-based methods (2nd order):
 - Uses information about gradients and curvature
 - Converges more rapidly to the solution
 - High cost of computing and storing Hessians
- Derivative-free methods (0th order):
 - Slow
 - Suitable for problems with kinks and discontinuities
- Trade-off: Higher order methods are speedier, while lower order methods give more accurate solutions

Several Common Methods

- Derivative-based methods
 - Bisection method (1st order)
 - Newton's method (2nd order)
 - Quasi-Newton method (2nd order)
- Derivative-free methods
 - Grid search method
 - Bracket method
 - Golden section search method
 - Nelder-Mead method
- Constrained optimization
 - Penalty function method
 - Sequential least squares programming (Derivative-based)
 - By linear approximation (Derivative-free)
- SciPy optimization package ([tutorial](#))

Derivative-Based Methods

Bisection Method (1st order)

- Similar logic to the bisection in non-linear equation solving, based on the Intermediate Value Theorem

Step 0 Find x^L & x^R s.t. $f'(x^L)f'(x^R) < 0$
Choose stopping criterion²: ϵ or δ

Step 1 Compute midpoint: $x^M = (x^L + x^R)/2$

Step 2 Update $x^L = x^L$ & $x^R = x^M$ if $f'(x^L)f'(x^M) < 0$
Update $x^L = x^M$ & $x^R = x^R$ if $f'(x^L)f'(x^M) > 0$

Step 3 Check stopping criterion:
If $x^R - x^L \leq \epsilon(1 + |x^L| + |x^R|)$ or $|f'(x^M)| \leq \delta$,
stop and report the solution at x^M
Otherwise, go back to Step 1 again

Newton's Method (2nd order)

- Similar logic to Newton's method in non-linear equation solving
- $f(x) \approx f(x_k) + f'(x_k)(x - x_k) + \frac{f''(x_k)}{2}(x - x_k)^2$
- FOC: $0 = f'(x^*) \approx f'(x_k) + (x^* - x_k)f''(x_k)$

Step 0 Set x_k with $k = 0$

Choose stopping criteria: ϵ & δ

Step 1 Compute $x_{k+1} = x_k - \frac{f'(x_k)}{f''(x_k)}$

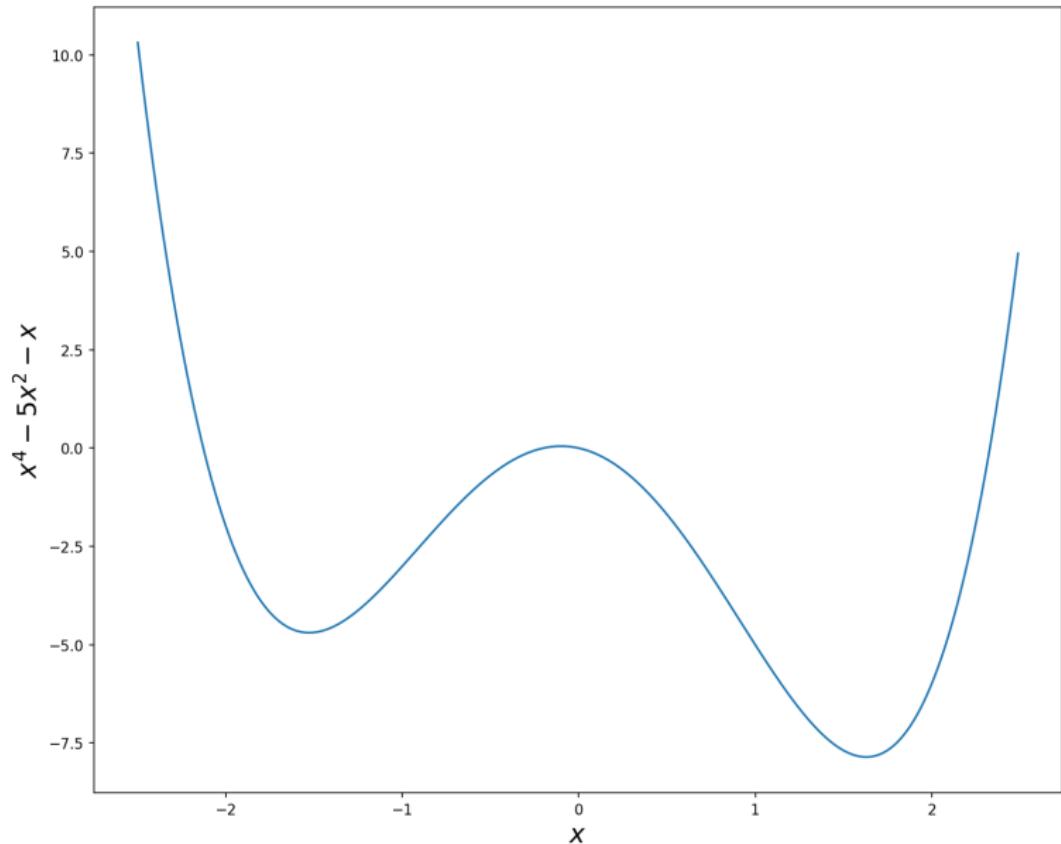
Step 2 Check stopping criterion:

If $|x_k - x_{k+1}| \leq \epsilon(1 + |x_{k+1}|)$, go to Step 3.

Otherwise, go back to Step 1.

Step 3 If $|f'(x_{k+1})| \leq \delta$, report $x^* = x_{k+1}$ as a solution.
Otherwise, report failure.

Eg.1) Local & Global Optimum



Eg.1) Initial Guess Matters

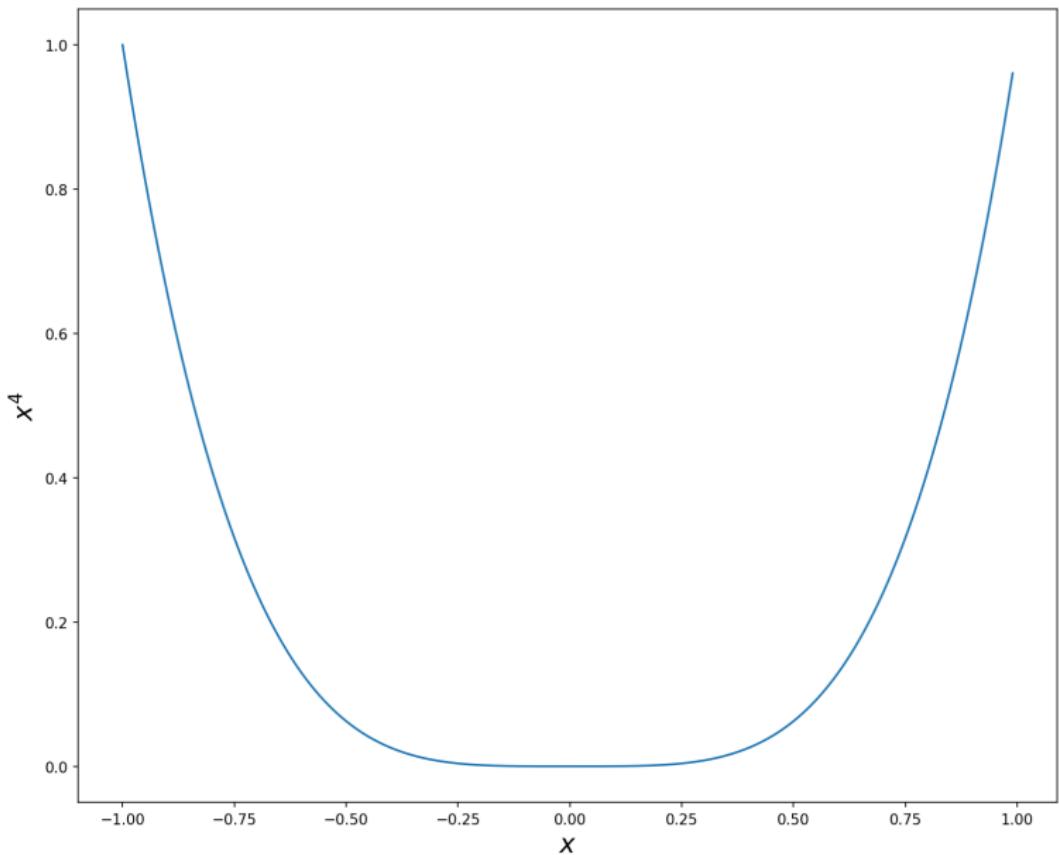
```
import matplotlib.pyplot as plt
import numpy as np
import scipy
from scipy.optimize import minimize

obj_1 = lambda x: x**4 - 5*(x**2) - x
argmin_1_1 = minimize(obj_1, x0=-0.5, method='BFGS') # a quasi-newton's method
argmin_1_2 = minimize(obj_1, x0=0.5, method='BFGS')

In [12]: argmin_1_1
Out[12]:
    fun: -4.694706337665813
    hess_inv: array([[ 0.05546626]])
    jac: array([-5.96046448e-08])
    message: 'Optimization terminated successfully.'
    nfev: 18
    nit: 4
    njev: 6
    status: 0
    success: True
    x: array([-1.52854364])

In [13]: argmin_1_2
Out[13]:
    fun: -7.855394472077334
    hess_inv: array([[ 0.0457116]])
    jac: array([-8.34465027e-07])
    message: 'Optimization terminated successfully.'
    nfev: 21
    nit: 5
    njev: 7
    status: 0
    success: True
    x: array([ 1.62894846])
```

Eg.2) Too Small Changes in Gradients



Eg.2) Scaling Matters

```
import matplotlib.pyplot as plt
import numpy as np
import scipy
from scipy.optimize import minimize

obj_2 = lambda x: x**4
obj_2_transform = lambda x: 10000*x**4

In [12]: minimize(obj_2, 10, method='BFGS')
Out[12]:
    fun: 1.1274113375014166e-08
    hess_inv: array([[ 328.87754176]])
    jac: array([ 4.37645767e-06])
    message: 'Optimization_terminated-successfully.'
    nfev: 78
    nit: 25
    njev: 26
    status: 0
    success: True
    x: array([ 0.01030435])

In [13]: minimize(obj_2_transform, 10, method='BFGS')
Out[13]:
    fun: 1.4693518450033662e-09
    hess_inv: array([[ 9.10967096]])
    jac: array([ 9.49336157e-06])
    message: 'Optimization_terminated-successfully.'
    nfev: 108
    nit: 35
    njev: 36
    status: 0
    success: True
    x: array([ 0.00061913])
```

Newton's Method (Multidimensional)

- $f : \mathbb{R}^n \rightarrow \mathbb{R}$; $H_f()$: Hessian of $f()$
 $\mathbf{0} = \nabla f(\mathbf{x}^*) \approx \nabla f(\mathbf{x}_k) + H_f(\mathbf{x}_k)(\mathbf{x}^* - \mathbf{x}_k)$

Step 0 Set \mathbf{x}_k with $k = 0$

Choose stopping criteria: ϵ & δ

Step 1 Compute $\mathbf{x}_{k+1} = \mathbf{x}_k - H_f(\mathbf{x}_k)^{-1} \nabla f(\mathbf{x}_k)$

Step 2 Check stopping criterion:

If $\|\mathbf{x}_k - \mathbf{x}_{k+1}\| \leq \epsilon(1 + \|\mathbf{x}_{k+1}\|)$, go to Step 3.

Otherwise, go back to Step 1.

Step 3 If $\|\nabla f(\mathbf{x}_k)\| \leq \delta(1 + |f(x_k)|)$, report $\mathbf{x}^* = \mathbf{x}_{k+1}$ as an optimum.
 Otherwise, report failure.

- Converge quadratically to a local optimum.

Newton's Methods (Multidimensional): Caveats

- Calculating Hessian and its inverse is costly

Newton's Methods (Multidimensional): Caveats

- Calculating Hessian and its inverse is costly
- Also, Hessian must be well-conditioned:
 - Invertible
 - Positive semi-definiteness (in a minimization problem) around the solution, so that the objective function value is approached toward the solution in each Newton step

Newton's Methods (Multidimensional): Caveats

- Calculating Hessian and its inverse is costly
- Also, Hessian must be well-conditioned:
 - Invertible
 - Positive semi-definiteness (in a minimization problem) around the solution, so that the objective function value is approached toward the solution in each Newton step
- However, no guarantee that the above conditions hold, especially at points far from the solution

Practical Methods for Multidimensional Cases

- **Quasi-Newton methods:** approximate a Hessian (or its inverse) by a positive definite H_k , guaranteeing that function value can be decreased in the direction of the Newton step (in a minimization problem), s.t.

$$H_k^{-1}(\nabla f(\mathbf{x}_{k+1})' - \nabla f(\mathbf{x}_k)') = \mathbf{x}_{k+1} - \mathbf{x}_k$$

Check several quasi-Newton methods with different updating rules of $\{H_k\}$ by yourself:

- Broyden-Fletcher-Goldfarb-Shanno (BFGS) method
- Davidson-Fletcher-Powell (DFP) method
- **Conjugate Gradient Method (CGM):** Store only a gradient, while implicitly keeping track of curvature information in a useful way without storing a Hessian
- See *Judd, Chapter 4* in detail

Derivative-Free Methods

Grid Search Method

- Simplest and most primitive

Grid Search Method

- Simplest and most primitive
- Credible global solution (if parameter space is sufficiently wide and grids are not too coarse)
- Slow (especially for a high dimensional case)

Grid Search Method

- Simplest and most primitive
- Credible global solution (if parameter space is sufficiently wide and grids are not too coarse)
- Slow (especially for a high dimensional case)
- Useful for understanding the shape of objective function:
 - Unless you are aware of the functional form clearly (which rarely happens), for whatever method you will finally adopt, begin by plotting with this method
 - That helps you to select which type of more sophisticated method to adopt if necessary

Nelder-Mead (Downhill Simplex) Method

- A widely-used derivative-free optimization method for multi-dimensional functions
- $f : \mathbb{R}^n \rightarrow \mathbb{R}$
- Begin by evaluating the objective function at $n + 1$ points
- These points form a *simplex* in \mathbb{R}^n
- Directly search for the optimum by moving this simplex with several steps
(reflection; expansion; contraction; shrinkage)

Nelder-Mead (Downhill Simplex) Method

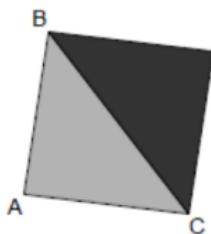
- A widely-used derivative-free optimization method for multi-dimensional functions
- $f : \mathbb{R}^n \rightarrow \mathbb{R}$
- Begin by evaluating the objective function at $n + 1$ points
- These points form a *simplex* in \mathbb{R}^n
- Directly search for the optimum by moving this simplex with several steps
(reflection; expansion; contraction; shrinkage)

Remark. Speedier than the grid search method (and much more accurate than derivative-based methods), but not still perfect for converging to a global solution. Try with several initial guesses.

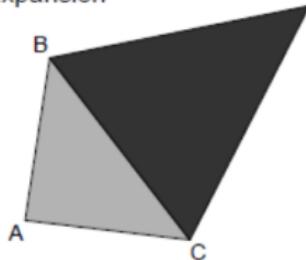
Illustration with $n = 2$

Simplex Transformations in the Nelder–Mead Algorithm

Reflection



Expansion



Contraction

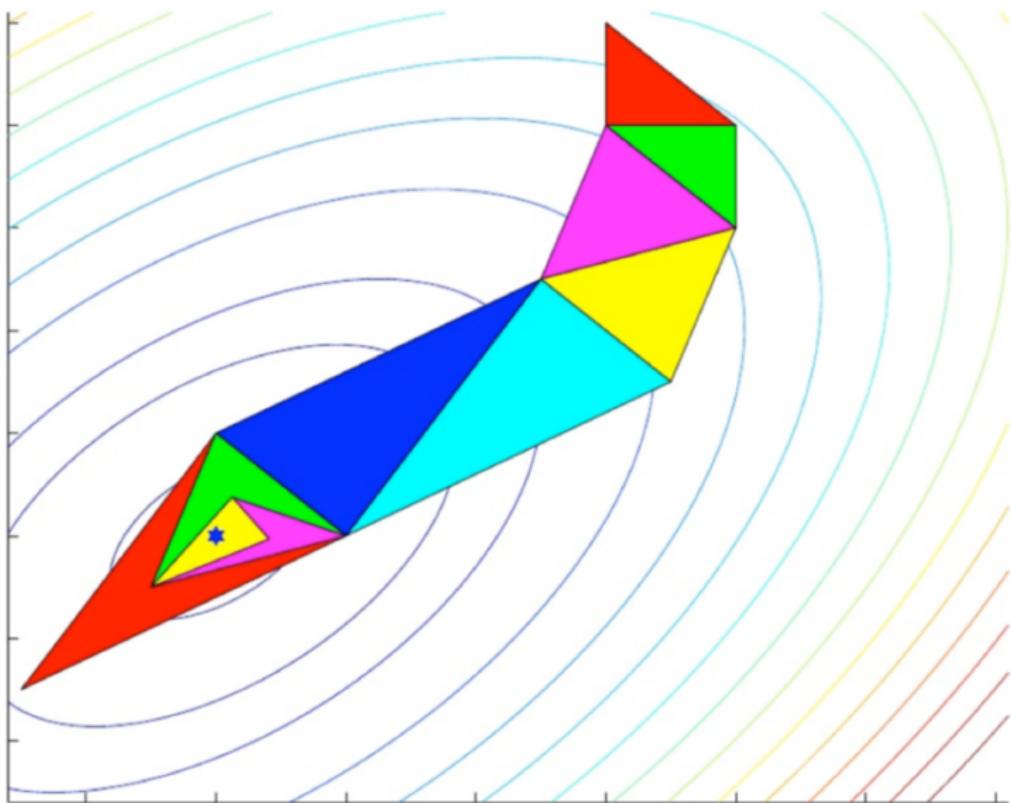


Shrinkage



(Miranda and Fackler, Chapter 4)

Illustration with $n = 2$



(Lecture note by Fernández-Villaverde and Guerrón)

Derivative-Based vs. Derivative-Free Methods

- Plot the objective function with some coarse grids to get a sense.
- If you are sure that the function is globally concave/convex, use a derivative-based method.
(Check if different initial guesses actually globally converge.)
- Otherwise, go for a derivative-free method like Nelder-Mead.
(Again, check with different initial guesses.)
- Whatever method you finally adopt, check the robustness with the grid search method in some reasonable range of parameters.

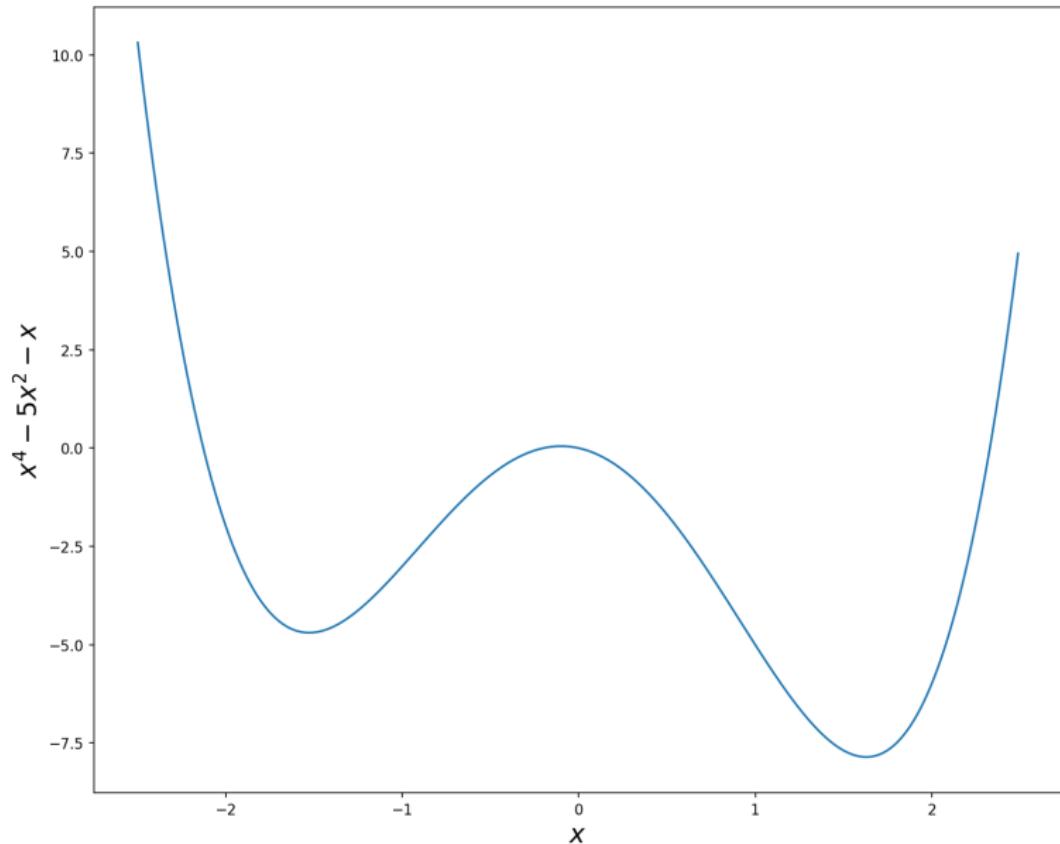
Derivative-Based vs. Derivative-Free Methods

- Plot the objective function with some coarse grids to get a sense.
- If you are sure that the function is globally concave/convex, use a derivative-based method.
(Check if different initial guesses actually globally converge.)
- Otherwise, go for a derivative-free method like Nelder-Mead.
(Again, check with different initial guesses.)
- Whatever method you finally adopt, check the robustness with the grid search method in some reasonable range of parameters.
- Trade-off: In general, higher order methods (derivative-based) are speedier, while lower order methods (derivative-free) methods tend to give more accurate solutions

Derivative-Based vs. Derivative-Free Methods

- Plot the objective function with some coarse grids to get a sense.
- If you are sure that the function is globally concave/convex, use a derivative-based method.
(Check if different initial guesses actually globally converge.)
- Otherwise, go for a derivative-free method like Nelder-Mead.
(Again, check with different initial guesses.)
- Whatever method you finally adopt, check the robustness with the grid search method in some reasonable range of parameters.
- Trade-off: In general, higher order methods (derivative-based) are speedier, while lower order methods (derivative-free) methods tend to give more accurate solutions
- Accuracy of derivative-free methods is not still perfect. (Eg. 1)

Eg.1) (Again) Local and global optimum



Eg.1) Newton's vs. Nelder-Mead

```
In [12]: minimize(obj_1, -0.1, method='BFGS')
Out[12]:
    fun: -7.855394472077359
    hess_inv: array([[ 419.3338208]])
    jac: array([ 1.19209290e-07])
    message: 'Optimization_terminated_successfully.'
    nfev: 36
    nit: 2
    njev: 12
    status: 0
    success: True
    x: array([ 1.6289485])

In [13]: minimize(obj_1, -0.1, method='Nelder-Mead')
Out[13]:
    final_simplex: (array([[-1.52851563],
                           [-1.52859375]]), array([-4.69470633, -4.69470632]))
    fun: -4.6947063305936503
    message: 'Optimization_terminated_successfully.'
    nfev: 42
    nit: 21
    status: 0
    success: True
    x: array([-1.52851563])
```

Derivative-Based vs. Derivative-Free Methods

- Plot the objective function with some coarse grids to get a sense.
- If you are sure that the function is globally concave/convex, use a derivative-based method.
(Check if different initial guesses actually globally converge.)
- Otherwise, go for a derivative-free method like Nelder-Mead.
(Again, check with different initial guesses.)
- Whatever method you finally adopt, check the robustness with the grid search method in some reasonable range of parameters.
- Trade-off: In general, higher order methods (derivative-based) are speedier, while lower order methods (derivative-free) methods tend to give more accurate solutions
- Accuracy of derivative-free methods is NOT still perfect. (Eg. 1)
- Derivative-free methods are superior especially when there are discontinuities in the objective function. In such cases, derivative-based methods work very poorly. (Eg. 3)

Eg.3) Discontinuities in the Objective Function

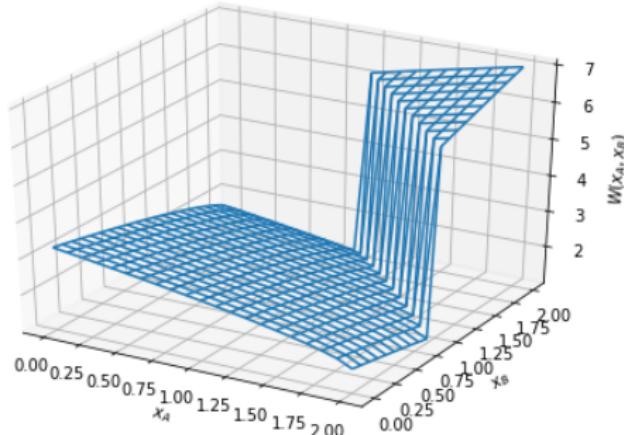
- 2-player public investment game

- Utility:

$$u_i(x_A, x_B) = \left(2 - x_i + \mathbb{1}_{x_A+x_B \geq 3} \frac{3(x_A+x_B)}{2}\right)^\sigma, i = A, B$$

- Welfare: $W(x_A, x_B) \equiv u_A(x_A, x_B) + u_B(x_A, x_B)$

- Welfare-maximizer: $(x_A^*, x_B^*) = (2, 2)$



Eg.3) Discontinuities in the Objective Function

- Utility:

$$u_i(x_A, x_B) = \left(2 - x_i + \mathbb{1}_{x_A+x_B \geq 3} \frac{3(x_A+x_B)}{2}\right)^\sigma, i = A, B$$

- Welfare: $W(x_A, x_B) \equiv u_A(x_A, x_B) + u_B(x_A, x_B)$

```
import numpy as np
from scipy.optimize import minimize
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D

def Inv_W(x):
    x_a = x[0]
    x_b = x[1]
    sigma = 0.7
    inv_return = np.where(x_a+x_b>=3, 3*(x_a+x_b), 0)
    return -((2-x_a) + inv_return/2)**sigma + ((2-x_b) + inv_return/2)**sigma

x_a = np.arange(0, 2.0+0.1, 0.1)
x_b = np.arange(0, 2.0+0.1, 0.1)
X_a, X_b = np.meshgrid(x_a, x_b)
PG_welfare = -Inv_W(np.array([X_a, X_b]))

ax = Axes3D(plt.figure())
ax.set_xlabel("$x_A$")
ax.set_ylabel("$x_B$")
ax.set_zlabel("$W(x_A, x_B)$")
ax.plot_wireframe(X_a, X_b, PG_welfare)
plt.show()
```

Set up for a constrained optimization problem:

```
constraint1 = lambda x: x-0.0
constraint2 = lambda x: 2.0-x
const1 = ({'type': 'ineq', 'fun' : constraint1})
const2 = ({'type': 'ineq', 'fun' : constraint2})
const = [const1, const2]
```

Derivative-based (Sequential Least Squares Programming):

```
In[12]: inv_init = [1.49, 1.49]
...: minimize(Inv_W, inv_init, method='SLSQP', constraints=const)
Out[12]:
    fun: -3.2490095854249414
    jac: array([0.56857666, 0.56857666])
message: 'Optimization_terminated_successfully.'
    nfev: 12
    nit: 3
    njev: 3
status: 0
success: True
    x: array([4.4408921e-16, 0.0000000e+00])
```

Derivative-free (Constrained Optimization By Linear Approximation):

```
In[13]: inv_init = [1.1, 1.1]
...: minimize(Inv_W, inv_init, method='COBYLA', constraints=const)
Out[13]:
    fun: -7.010346004639045
    maxcv: 3.535533905951738e-05
message: 'Optimization_terminated_successfully.'
    nfev: 13
    status: 1
success: True
    x: array([2.00003536, 2.00003536])
```

Practical Advice: Solving

- First of all, judge if numerical optimization is really necessary

Practical Advice: Solving

- First of all, judge if numerical optimization is really necessary
- Estimate the objective function with multiple optimization algorithms, ideally from different orders
- Try multiple stopping rules based on tight tolerances
- Try multiple initial guesses for parameter values:
 - Based on prior information from economic theory
 - Without such information, use random draws

Practical Advice: Solving

- First of all, judge if numerical optimization is really necessary
- Estimate the objective function with multiple optimization algorithms, ideally from different orders
- Try multiple stopping rules based on tight tolerances
- Try multiple initial guesses for parameter values:
 - Based on prior information from economic theory
 - Without such information, use random draws
- (For a very high dimensional problem,) divide the optimization problem in hand in smaller dimensions:
 - Split the parameter vector θ into θ_1 & θ_2
 - Fix θ_1 and estimate θ_2
 - With estimates of θ_2 , estimate θ_1
 - Iterate this process until the objective function value changes little
 - Finally, simultaneously estimate θ_1 & θ_2

Practical Advice: Solving

- First of all, judge if numerical optimization is really necessary
- Estimate the objective function with multiple optimization algorithms, ideally from different orders
- Try multiple stopping rules based on tight tolerances
- Try multiple initial guesses for parameter values:
 - Based on prior information from economic theory
 - Without such information, use random draws
- (For a very high dimensional problem,) divide the optimization problem in hand in smaller dimensions:
 - Split the parameter vector θ into θ_1 & θ_2
 - Fix θ_1 and estimate θ_2
 - With estimates of θ_2 , estimate θ_1
 - Iterate this process until the objective function value changes little
 - Finally, simultaneously estimate θ_1 & θ_2
- Consider combining optimization algorithms:
 - Grid search with coarse grids for guessing a reasonable range of parameter values
 - Then, use other faster methods to estimate the parameters

Practical Advice: Validating

- Collect the sets of parameter estimates & objective function values
- Pick the parameter estimates which achieve the lowest objective function value among the collected sets
- Make sure that the algorithm converges to the picked parameter estimates

Practical Advice: Validating

- Collect the sets of parameter estimates & objective function values
 - Pick the parameter estimates which achieve the lowest objective function value among the collected sets
 - Make sure that the algorithm converges to the picked parameter estimates
 - Plot the objective function value against each of the parameters over an interval around the estimated value, keeping the other parameters fixed at the estimated values
- ⇒ Make sure it passes the “ocular” identification test

Practical Advice: Validating

- Collect the sets of parameter estimates & objective function values
 - Pick the parameter estimates which achieve the lowest objective function value among the collected sets
 - Make sure that the algorithm converges to the picked parameter estimates
 - Plot the objective function value against each of the parameters over an interval around the estimated value, keeping the other parameters fixed at the estimated values
- ⇒ Make sure it passes the “ocular” identification test
- Simulate a dataset given the estimated parameter values
 - ⇒ Estimate the parameters given the simulated dataset and make sure that the newly estimated parameters have the same values as the ones obtained first

Practical Advice: Validating

- Collect the sets of parameter estimates & objective function values
 - Pick the parameter estimates which achieve the lowest objective function value among the collected sets
 - Make sure that the algorithm converges to the picked parameter estimates
 - Plot the objective function value against each of the parameters over an interval around the estimated value, keeping the other parameters fixed at the estimated values
- ⇒ Make sure it passes the “ocular” identification test
- Simulate a dataset given the estimated parameter values
 - ⇒ Estimate the parameters given the simulated dataset and make sure that the newly estimated parameters have the same values as the ones obtained first
 - **Again, no perfect procedure for identifying global optima exists!**
But think wisely for reducing the possibilities of failures!

TABLE 4.—OPTIMIZATION-DESIGN DETAILS AND DIAGNOSTICS CHECKLIST

Step	Details and Diagnostics
1. Optimization design	Optimization algorithm Starting values Objective function value tolerance Parameter vector tolerance Other optimization settings Fixed-point iteration settings Market share evaluation draws
2. Convergence and local optima	Multiple optima (Y/N) Number of runs converged Algorithm exit code Objective function value Parameter estimates Gradient-based FOC diagnostics Hessian-based SOC diagnostics
3. Implications for economic variables of interest	Variation due to multiple optima, if any, for: <ul style="list-style-type: none"> • Objective function value • Parameter estimates • Own- and cross-price elasticities: statistics • Other economic variables of interest: statistics

Source: Nittel & Metaxoglou (2014)

Further Readings

There are still much more methods that this lecture has not covered.
Please see, for example:

- *Judd, Chapter 4*
- Note by Fernández-Villaverde & Guerrón
- Note by Todd Munson
- Knittel, Christopher R., and Konstantinos Metaxoglou.
“Estimation of random-coefficient demand models: two empiricists’ perspective.” *Review of Economics and Statistics* 96.1 (2014): 34-59.
- Also, read SciPy documentations carefully

Lab

Exercise

Aim. Experience the imperfection of numerical optimization for finding a global solution

Task. Minimize $f(x) = 3x^4 - 5x^3 + 2x^2$ by

1. (Quasi-)Newton's and Nelder-Mead methods with initial guesses

1.1. -0.25

1.2. 0

1.3. 0.25

1.4. 0.5

1.5. 0.75

1.6. 1

2. Grid search

Assignment 4

Setting: Quasi-hyperbolic discounting structure

- An application to a simple structural model
- Time preferences play important roles for various dynamic decisions.
- Not only discount factor, but also present biasness matters.
e.g.) I do not want to do my homework just now, so I allocate much more study time to tomorrow. The ratio of study time between today and a future day can differ from the planned ratio between two future days with an equal interval.
- An individual at period t maximizes lifetime utility:

$$U(c) = u(c_t) + \beta \sum_{k=1}^{\infty} \delta^k u(c_{t+k})$$

- Read Andreoni and Sprenger (2012 AER), Augenblick et al. (2015 QJE), and Casaburi and Macchiavello (2019 AER) if you are interested, but not necessary for this assignment.

Experiment & Data Generating Process

- A researcher conducts a lab experiment in India for obtaining time preference parameters.
- Subjects are asked to choose two-period intertemporal allocations of money (Rs. 4000) within a convex budget set, with various t (earlier date), k (time interval between the earlier and later dates), and several interest rates $P = (1 + r)$.
- Assume that the subjects solve:

$$\begin{aligned} U(c_t, c_{t+k}) &= c_t^\alpha + \beta^{\mathbb{1}_{t=0}} \delta^k c_{t+k}^\alpha \\ \text{s.t. } &Pc_t + c_{t+k} = 4000 \end{aligned}$$

- Solving this,

$$c_t = \frac{4000(\beta^{\mathbb{1}_{t=0}} \delta^k P)^{1/(\alpha-1)}}{1 + P(\beta^{\mathbb{1}_{t=0}} \delta^k P)^{1/(\alpha-1)}} \equiv g(\mathbb{1}_{t=0}, k, P; \beta, \delta, \alpha)$$

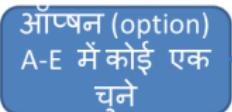
- Parameters: β : present biasness, δ : discount factor, α : curvature ($IES = 1/(1 - \alpha)$).

Context Time Budget (CTB) Experiment

Section C: निर्देश

- Test C-1 के answer sheet का एक उदाहरण

आँप्षन (option) A-E में कोई एक चुने



	Option A	Option B	Option C	Option D	Option E
Today	3800 Rs. & 0 Rs.	2850 Rs. & 1000 Rs.	1900 Rs. & 2000 Rs.	950 Rs. & 3000 Rs.	0 Rs. & 4000 Rs.
5 weeks from today					

Context Time Budget (CTB) Experiment (cont'd)

Answer sheet for Test C-19

30

	Option A	Option B	Option C	Option D	Option E
9 weeks later →	<input type="checkbox"/>				
	4000 Rs.	3000 Rs.	2000 Rs.	1000 Rs.	0 Rs.
	&	&	&	&	&
18 weeks later →	0 Rs.	1000 Rs.	2000 Rs.	3000 Rs.	4000 Rs.

Assignment: Non-Linear Least Square Estimation

Aim. Validate the numerical solution to a nonlinear problem

- Distributed data: $w_{i,q} \equiv \{c_{i,t_q}, c_{i,t_q+k_q}, t_q, k_q, P_q\}$
(i : individual, q : question)

Task. Estimate the parameters $\hat{\theta} = (\hat{\beta}, \hat{\delta}, \hat{\alpha})$ by non-linear least squares (NLLS):

$$\min_{\beta, \delta, \alpha} \sum_{i,q} [c_{i,t_q} - g(\mathbf{1}_{t_q=0}, k_q, P_q; \beta, \delta, \alpha)]^2$$

- Try both derivative-based and derivative-free methods.
- Feel free to use the SciPy optimization package.
- Follow the practical advice to validate your result.
E.g.) Plot each parameter value and the objective function, fixing the other parameter values at the estimated values.